



[Unit 4 Unsupervised Learning \(2](#)

[Course](#) > [weeks](#))

> [Lecture 15. Generative Models](#) >

6. MLE for Multinomial Distribution

**Audit Access Expires May 11, 2020**

You lose all access to this course, including your progress, on May 11, 2020.

## 6. MLE for Multinomial Distribution

### Maximum Likelihood Estimate



#### Video

[Download video file](#)

#### Transcripts

[Download SubRip \(.srt\) file](#)

[Download Text \(.txt\) file](#)

## Deriving MLE for a General Multinomial Model: Likelihood

1/1 point (graded)

In the following problems, we will derive the maximum likelihood estimates for a multinomial model with more than 2 parameters. We will employ the method of lagrange multipliers for the optimization problem.

Let the document  $D$  be a sequence of words  $w_1, \dots, w_n$  from a collection  $W$  consisting of  $N$  words. For simplicity, we assume that  $w_i$ 's are independent, and that the probability of a word  $w$  is given by the parameter  $\theta_w$ , and denote by  $\theta = \{\theta_w\}_{w \in W}$ .

Let  $P(D|\theta)$  be the probability of  $D$  being generated by the simple model described above.

Find  $P(D|\theta)$ .

☐  $P(D|\theta) = \sum_{w \in W} \theta_w^{\text{count}(w)}$

☒  $P(D|\theta) = \prod_{w \in W} \theta_w^{\text{count}(w)}$

☐  $P(D|\theta) = \prod_{w \in W} \text{count}(w)^{\theta_w}$

☐  $P(D|\theta) = \prod_{w \in W} \theta_w + \text{count}(w)$



**Solution:**

The probability of  $D$

$$P(D|\theta) = \prod_{i=1}^n \theta_{w_i} = \prod_{w \in W} \theta_w^{\text{count}(w)}$$

**Submit**

You have used 1 of 2 attempts

 Answers are displayed within the problem

## Constraints on the Parameters

1/1 point (graded)

What are the constraints on the parameters  $\theta_w$  in the model described in the previous problem?

☒  $\theta_w \geq 0, \sum_{w \in W} \theta_w = 1$

☐  $\theta_w \geq 0, \sum_{w \in W} \theta_w < 1$

☐  $\theta_w < 0, \sum_{w \in W} \theta_w > -1$

☐  $\theta_w \geq 0, \sum_{w \in W} \theta_w \geq 1$



### Solution:

Since  $\theta_w$  denotes the probability of the word  $w$  under the model, its value must lie between 0 and 1. Therefore,  $0 \leq \theta_w \leq 1$ .

Further, all the above probability values must also sum up to 1. That is,  $\sum_{w \in W} \theta_w = 1$ .

**Submit**

You have used 1 of 2 attempts

 Answers are displayed within the problem

## Stationary Points of the Lagrange Function

2/2 points (graded)

The maximum likelihood estimate of  $\theta$  is the value of  $\theta$  that maximizes the likelihood function:

$$P(D|\theta) = \prod_{w \in W} (\theta_w)^{\text{count}(w)}.$$

Maximizing  $P(D|\theta)$  is equivalent to maximizing  $\log P(D|\theta)$ , so we take the natural logarithm on both sides of the equation above to bring down the exponents:

$$\log P(D|\theta) = \sum_{w \in W} \text{count}(w) \log \theta_w.$$

Recall that  $\theta$  is subject to the following constraint:

$$\sum_{w \in W} \theta_w = 1.$$

To maximize  $\log P(D|\theta)$  subject to the constraint  $\sum_{w \in W} \theta_w = 1$ , we use the Lagrange multiplier method.

### Method of Lagrange Multipliers

#### Problem

Let  $f$  be a function from  $\mathbb{R}^N$  to  $\mathbb{R}$ . Find the (local) maxima/minima of  $f$  **subject to a given constraint**  $g = 0$ , where  $g$  is a function  $\mathbb{R}^N$  to  $\mathbb{R}$ .

A two dimensional example is: Find the local extrema of  $f(x, y) = x^2$  subject to the constraint  $x^2 + y^2 = 1$  i.e. optimize the function  $f$  on the unit circle.

#### Method of Lagrange Multipliers

Without the constraint, the optimization problem can be solved as usual by setting the gradient of  $f$  to zero i.e.

$$\nabla f = 0.$$

With the constraint, we can solve the following equation instead:

$$\nabla f = \lambda \nabla g$$

where  $\lambda$  is a constant scalar. Geometrically, for  $\lambda \neq 0$ , a solution to the equation above is a point in  $\mathbb{R}^N$  where the gradient of  $f$  is "parallel" to the gradient of  $g$ , or equivalently, where the gradient of  $f$  is perpendicular to the tangent of the curve defined by  $g = k$  for some  $k$ . In other words, at a solution point, the directional derivative of  $f$  is zero along the direction tangent to the curve  $g = k$  for some constant  $k$ , and hence  $f$  is stationary as we travel along  $g = k$ .

Finally, we impose the constraint  $g = 0$  to find the local extrema of  $f$  on  $g = 0$ .

Since the equation  $\nabla f = \lambda \nabla g$  is equivalent to  $\nabla L = 0$  where  $L = f - \lambda g$ , the problem of optimizing  $f$  subject to  $g = 0$  can be reformulated as optimizing the function  $L$  along with the constraint  $g = 0$ . We call the function  $L$  the **Lagrangian function**, and the scalar  $\lambda$  the **Lagrange multiplier**.

Note that we can equally define  $L = f + \lambda g$ , since  $\lambda$  is an unknown scalar we will solve for.

### Example

Find the local extrema of  $f(x, y) = x^2$  subject to the constraint  $x^2 + y^2 = 1$ . Geometrically, the function  $f$  is a parabolic cylinder, i.e.  $f$  is a parabolic in the  $x$  direction with constant values in the  $y$  direction. The constraint is a unit circle.

Solution:

First, solve the equation

$$\nabla f = \lambda \nabla g \quad \text{where } g(x, y) = x^2 + y^2 - 1$$

$$\begin{aligned} \iff \begin{bmatrix} 2x \\ 0 \end{bmatrix} &= \lambda \begin{bmatrix} 2x \\ 2y \end{bmatrix} \\ \iff \begin{bmatrix} (1 - \lambda) 2x \\ \lambda (2y) \end{bmatrix} &= 0 \end{aligned}$$

The set of all possible solutions to the equation above are  $(\lambda = 1, y = 0)$ , or  $(\lambda = 0, x = 0)$ , or  $(x = y = 0)$ .

Finally, impose the constraint  $x^2 + y^2 - 1 = 0$  to further pin down the local extrema. Subject to  $x^2 + y^2 = 1$ ,  $f(x, y) = x^2$  is at local maximum or minimum at  $(x = 0, y = \pm 1)$  and  $(y = 0, x = \pm 1)$ . At  $(x = 0, y = \pm 1)$ , we have  $\lambda = 0$  and  $\nabla f = 0$ . Since  $f$  has only local minima, these two points remain local minima of  $f$  on the unit circle. At  $(y = 0, x = \pm 1)$ , we have  $\lambda = 1$  and hence  $\nabla f = \nabla g$ . Equivalently, the directional derivative  $\nabla f$  is zero along the tangent direction of the circle at this point. Visualizing or computing second derivatives will allow us to see that these two points are local maxima of  $f$  along the unit circle.

[Hide](#)

Define the Lagrange function:

$$L = \log P(D|\theta) + \lambda \left( \sum_{w \in W} \theta_w - 1 \right)$$

where  $\lambda$  is a constant scalar.

Then, find the stationary points of  $L$  by solving the equation  $\nabla_{\theta} L = 0$ . The components of this equation are

$$\frac{\partial}{\partial \theta_w} \left( \log P(D|\theta) + \lambda \left( \sum_{w \in W} \theta_w - 1 \right) \right) = 0 \quad \text{for all } w \in W.$$

Solve for  $\theta_w$  from the above equation. Choose the right answer for  $\theta_w$  from options below.

☐  $\theta_w = \frac{-\lambda}{\text{count}(w)}$

☐  $\theta_w = \lambda \text{count}(w)$

☐  $\theta_w = -\lambda \text{count}(w)$

☒  $\theta_w = \frac{-\text{count}(w)}{\lambda}$



Now, apply the constraint that  $\sum_{w \in W} \theta_w = 1$  to the answer above to obtain  $\lambda$ .

$\lambda =$

☒  $\lambda = - \sum_{w \in W} \text{count}(w)$

☐  $\lambda = \sum_{w \in W} \text{count}(w)$

☐  $\lambda = - \sum_{w \in W} (\theta_w \text{count}(w))$

☐  $\lambda = \sum_{w \in W} (\theta_w \text{count}(w))$



Find  $\theta_w$  that maximizes  $\log P(D|\theta)$  subject to  $\sum_{w \in W} \theta_w = 1$ .

(There is no answer box for this final question.)

**Solution:**

$$\frac{\partial}{\partial \theta_w} (\log P(D|\theta) + \lambda (\sum_{w \in W} \theta_w - 1)) = 0$$

$$\frac{\partial \log P(D|\theta_w)}{\partial \theta_w} + \lambda = 0$$

$$\frac{\partial \log \prod_{w \in W} \theta_w^{\text{count}(w)}}{\partial \theta_w} + \lambda = 0$$

$$\frac{\partial \sum_{w \in W} \log \theta_w \times \text{count}(w)}{\partial \theta_w} + \lambda = 0$$

$$\frac{\text{count}(w)}{\theta_w} + \lambda = 0$$

$$\theta_w = -\frac{\text{count}(w)}{\lambda}$$

If we apply the constraint that  $\sum_{w \in W} \theta_w = 1$  we get



$$\sum_{w \in W} \theta_w = 1$$

$$\sum_{w \in W} -\frac{\text{count}(w)}{\lambda} = 1$$

$$\sum_{w \in W} \text{count}(w) = -\lambda$$

$$\lambda = -\sum_{w \in W} \text{count}(w)$$

Substituting this expression for  $\lambda$  back into our previous expression for  $\theta_w$  we get

$$\theta_w = -\frac{\text{count}(w)}{\lambda}$$

$$\theta_w = \frac{\text{count}(w)}{\sum_{w \in W} \text{count}(w)}$$

Note that  $\theta_w > 0$  and  $\sum_{w \in W} \theta_w = 1$  satisfying all the constraints. These set of  $\theta_w$  parameters are the maximum likelihood estimates for this multinomial generative distribution.

Submit

You have used 1 of 3 attempts



 Answers are displayed within the problem

Discussion

Hide Discussion

**Topic:** Unit 4 Unsupervised Learning (2 weeks) :Lecture 15.  
Generative Models / 6. MLE for Multinomial Distribution

Add a Post

Show all posts	by recent activity
 <a href="#">Lagrange multiplier missing details?</a>	6
 <a href="#">For those who don't understand lagrange method from the text</a> <a href="#">Hope this video would help https://www.youtube.com/watch?v=yuqB-d5MjZA</a>	1

© All Rights Reserved