Audit Access Expires May 11, 2020
You lose all access to this course, including your progress, on May 11, 2020.

2. Learning to Control: Introduction to Reinforcement Learning
Learning to Control: Introduction Reinforcement Learning (RL)

Video
Download video file

Transcripts
Download SubRip (.srt) file
Download Text (.txt) file

Objectives of RL

1/1 point (graded)
What is the goal of reinforcement learning (RL), and how does it work?
(Choose all that apply.)

- [ ] The goal of RL is to minimize the loss between predictions and labels for points in a labelled training dataset
- [x] The goal of RL is to learn a good policy with none or limited supervision
- [ ] For an RL algorithm to work, it must a receive a non-zero reward after every step
- [ ] RL works by maximizing the reward for each immediate next step

✔

Solution:

- Unlike with Supervised Learning, there would typically be no labelled training dataset associated with Reinforcement Learning tasks. RL algorithms learn to pick “good" actions based on the rewards that they receive during training.
- RL is most applicable to tasks where there is no clear cut supervised training data available.

Reinforcement learning algorithms can learn to take actions so as to maximize some notion of a cumulative reward instead of the reward for the next step and they can take “good" actions even without any intermediate rewards.

Submit You have used 1 of 2 attempts

Answers are displayed within the problem

Discussion

Hide Discussion

Topic: Unit 5 Reinforcement Learning (2 weeks) :Lecture 17. Reinforcement Learning 1 / 2. Learning to Control: Introduction to Reinforcement Learning

Add a Post

Show all posts by recent activity

- True or False
  I did not understand why 3 and 4 are false? 4
- How RL works
  According to the lecture, isn't RL about maximising the cumulative reward at the end? 2