Audit Access Expires May 11, 2020
You lose all access to this course, including your progress, on May 11, 2020.

3. Q value iteration by sampling
Q value iteration by sampling

0:00 / 0:00 1.25x

Video
Download video file

Transcripts
Download SubRip (.srt) file
Download Text (.txt) file

Let us consider a toy example which might not be very realistic but which neverthless can help delineate the Q-value iteration for RL using sampling approach.

For this example, assume that there are only two states, s_1, s_2 and only one action possible from each of these states. Let a_{s_1}, a_{s_2} be the actions that could be taken from s_1 and s_2 respectively.

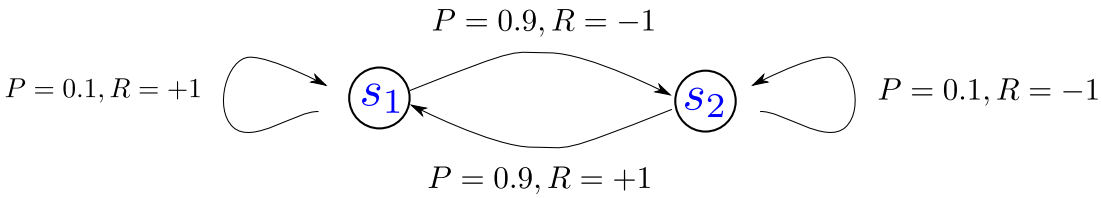

The state transition probabilities are listed below and are also shown in the figure above.

$$T(s_1, a_{s_1}, s_1) = 0.1$$

$$T(s_1, a_{s_1}, s_2) = 0.9$$

$$T(s_2, a_{s_2}, s_2) = 0.1$$

$$T(s_2, a_{s_2}, s_1) = 0.9$$

The rewards for these actions are given by

$$R(s_1, a_{s_1}, s_1) = 1$$

$$R(s_1, a_{s_1}, s_2) = -1$$

$$R(s_2, a_{s_2}, s_2) = -1$$

$$R(s_2, a_{s_2}, s_1) = 1$$

Note that we resort to finding optimal Q^* function by sampling for tasks where we don't have access to the exact T, R functions. However, for this toy example we will assume that the Q-value iteration algorithm isn't directly provided with the above specified values of T, R and has to resort to sampling to estimate the Q function.

Let's say that the agent starts out from state s_1 and collects few samples. Each sample can be described by the following tuple $(s, a, s', R(s, a, s'))$ which indicates that the agent received a reward of $R(s, a, s')$ when it reached state s' by taking action a from the state s.

The collected samples are described as follows in the order in which they are presented to the Q-value iteration algorithm.

$$(s_1, a_{s_1}, s_1, +1)$$

$$(s_1, a_{s_1}, s_2, -1)$$

$$(s_2, a_{s_2}, s_1, +1)$$

Let $S_k^{Q(s,a)}$ be used to denote the k^{th} sample of $Q(s, a)$ ($k = i + 1$). Then recall that

$$\hat{Q}_{i+1}(s, a) = \alpha * S_k^{Q(s,a)} + (1 - \alpha) * \hat{Q}_i(s, a)$$

For all of the following problems, assume that the discount factor $\gamma = 0.5$, $\alpha = 0.75$ and that all the Q values are initialized to 0 to start with. That is,

$$\hat{Q}_0(s, a) = 0 \forall s, a$$

Numerical Example

1 point possible (graded)
Enter below the value of $Q(s_1, a_{s_1})$ after the first sample is processed by the Q-value iteration algorithm

Answer: 0.75

Solution:

Let $S_k^{Q(s,a)}$ be used to denote the k^{th} sample of $Q\left(s,a\right)$.

$$\begin{aligned} S_1^{Q(s_1,a_{s_1})} &= R\left(s_1,a_{s_1},s_1\right)+\gamma*\max_{a'}Q\left(s_1,a'\right) \\ S_1^{Q(s_1,a_{s_1})} &= +1+0.5*0=1 \\ Q_1\left(s_1,a_{s_1}\right) &= \alpha*S_1^{Q(s_1,a_{s_1})}+(1-\alpha)*Q_0\left(s_1,a_{s_1}\right) \\ Q_1\left(s_1,a_{s_1}\right) &= .75*1+(1-.75)*0=.75 \end{aligned}$$

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

Numerical Example - 2

1 point possible (graded)
Enter below the value of $Q\left(s_1,a_{s_1}\right)$ after the second sample is seen by the Q-value iteration algorithm

Answer: -0.5625

Solution:

Let $S_k^{Q(s,a)}$ be used to denote the k^{th} sample of $Q\left(s,a\right)$. Note that from the previous example,

$$Q_1\left(s_1,a_{s_1}\right)=0.75$$

Now we find $S_2^{Q(s_1,a_{s_1})}$:

$$S_2^{Q(s_1,a_{s_1})} = R(s_1, a_{s_1}, s_2) + \gamma * \max_{a'} Q(s_2, a')$$
$$S_2^{Q(s_1,a_{s_1})} = -1 + 0.5 * 0 = -1$$
$$Q_2(s_1, a_{s_1}) = \alpha * S_2^{Q(s_1,a_{s_1})} + (1-\alpha) * Q_1(s_1, a_{s_1})$$
$$Q_2(s_1, a_{s_1}) = 0.75 * -1 + 0.25 * 0.75 = -0.5625$$

Submit You have used 0 of 3 attempts

Answers are displayed within the problem

Discussion

Hide Discussion

Topic: Unit 5 Reinforcement Learning (2 weeks) :Lecture 18. Reinforcement Learning 2 / 3. Q value iteration by sampling

Add a Post

Show all posts by recent activity

- Can someone explain the maxQ — Community TA — 28
- Incorrect formula on board — The last equation on the board in the lecture is incorrect. She has left out a factor of alpha. I.... — 2
- Do STATES increase during RL? — I have a question regarding the many things that are unknown, apart from ***T*** and ***R... — 3
- Hint for Q1 and Q2 — Community TA — 3
- Heraclitus vs. Markov: Can one ever step in the same river twice? — Is anyone working on a Heraclitus Decision Process (HDP) model? — 2
- Reinforcement learning resources — Survey paper: https://arxiv.org/abs/cs/9605103 Free textbook: http://incompleteideas.net/bo... — 1

?	Expression in answers Hi all, should the answers to questions in this vertical be inserted as mathematical expressio...	5
💬	[Staff] Hand Writing Difficult To Read The hand writing in this particular video is difficult to follow, it is all squished together. Please...	3
💬	[Staff] Second question bugged grader I introduced an incorrect answer which was graded as correct.	2
?	[Staff] Lecture formula for Q different from exercise?	5
💬	Q value in Max over a' of Q(s_1,a') is the value of Q impacted by the Q initialization to 0?	2
☑	[STAFF] Not clear how an exponentially weighted average formula leads to the recursive version in the lecture	2