

EdX and its Members use cookies and other tracking technologies for performance, analytics, and marketing purposes. By using this website, you accept this use. Learn more about these technologies in the [Privacy Policy](#).



[Unit 5 Reinforcement Learning](#)(2

[Course](#) > [weeks](#))

> [Project 5: Text-Based Game](#) >

3. Q-learning Algorithm

### Audit Access Expires May 25, 2020

You lose all access to this course, including your progress, on May 25, 2020.

## 3. Q-learning Algorithm

In this section, you will implement the Q-learning algorithm, which is a model-free algorithm used to learn an optimal Q-function. In the tabular setting, the algorithm maintains the Q-value for all possible state-action pairs. Starting from a random Q-function, the agent continuously collects experiences  $(s, c, R(s, c), s')$  and updates its Q-function.

From now on, we will refer to  $c = (a, b)$  as "an action" although it really is an action with an object.

### Q-learning Algorithm

- The agent plays an action  $c$  at state  $s$ , getting a reward  $R(s, c)$  and observing the next state  $s'$ .
- Update the single Q-value corresponding to each such transition:

Generating Speech Output

$$Q(s, c) \leftarrow (1 - \alpha) Q(s, c) + \alpha [R(s, c) + \gamma \max_{c' \in C} Q(s', c')]$$

**Tip:** We recommend you implement all functions from this tab and the next one before submitting your code online. Make sure you achieve reasonable performance on the *Home World* game

## Single step update

1.0/1 point (graded)

Write a function `tabular_q_learning` that updates the single Q-value, aiven the transition date  $(s, c, R(s, c), s')$ .

**Reminder:** You should implement this function locally first. You can read through the next tab to understand the context in which this function is called

**Available Functions:** You have access to the NumPy python library as `np`. You should also use constants `ALPHA` and `GAMMA` in your code

```

1 def tabular_q_learning(q_func, current_state_1, current_state_2, action
2     object_index, reward, next_state_1, next_state_2
3     terminal):
4     """Update q_func for a given transition
5
6     Args:
7         q_func (np.ndarray): current Q-function
8         current_state_1, current_state_2 (int, int): two indices descri
9         action_index (int): index of the current action
10        object_index (int): index of the current object
11        reward (float): the immediate reward the agent recieves from pl
12        next_state_1, next_state_2 (int, int): two indices describing t
13        terminal (bool): True if this epsiode is over
14
15    Returns:
```

Press ESC then TAB or click outside of the code editor to exit

Correct

Generating Speech Output

```
def tabular_q_learning(q_func, current_state_1, current_state_2, action_index,
                      object_index, reward, next_state_1, next_state_2,
                      terminal):
    """Update q_func for a given transition

    Args:
        q_func (np.ndarray): current Q-function
        current_state_1, current_state_2 (int, int): two indices describing the
        action_index (int): index of the current action
        object_index (int): index of the current object
        reward (float): the immediate reward the agent receives from playing cu
        next_state_1, next_state_2 (int, int): two indices describing the next
        terminal (bool): True if this episode is over

    Returns:
        None
    """
    if terminal:
        maxq_next_state = 0
    else:
        q_values_next_state = q_func[next_state_1, next_state_2, :, :]
        maxq_next_state = np.max(q_values_next_state)

    q_value = q_func[current_state_1, current_state_2, action_index,
                     object_index]
    q_func[current_state_1, current_state_2, action_index, object_index] = (
        1 - ALPHA) * q_value + ALPHA * (reward + GAMMA * maxq_next_state)
```

## Test results

**CORRECT**

[See full output](#)

[See full output](#)

Submit

You have used 11 of 25 attempts

**i** Answers are displayed within the problem

Generating Speech Output

## Epsilon-greedy exploration

1.0/1 point (graded)

Note that the Q-learning algorithm does not specify how we should interact in the world so as to learn quickly. It merely updates the values based on the experience collected. If we explore randomly, i.e., always select actions at random, we would most likely not get anywhere. A better option is to exploit what we have already learned, as summarized by current Q-values. We can always act greedily with respect to the current estimates, i.e., take an action  $\pi(s) = \arg \max_{c \in C} Q(s, c)$ . Of course, early on, these are not necessarily very good actions. For this reason, a typical exploration strategy is to follow a so-called  $\varepsilon$ -greedy policy: with probability  $\varepsilon$  take a random action out of  $C$  with probability  $1 - \varepsilon$  follow  $\pi(s) = \arg \max_{c \in C} Q(s, c)$ . The value of  $\varepsilon$  here balances exploration vs exploitation. A large value of  $\varepsilon$  means exploring more (randomly), not using much of what we have learned. A small  $\varepsilon$ , on the other hand, will generate experience consistent with the current estimates of Q-values.

Now you will write a function `epsilon_greedy` that implements the  $\varepsilon$ -greedy exploration policy using the current Q-function.

**Reminder:** You should implement this function locally first. You can read through the next tab to understand the context in which this function is called

**Available Functions:** You have access to the NumPy python library as `np`. Your code should also use constants `NUM_ACTIONS` and `NUM_OBJECTS`.

```

1 def epsilon_greedy(state_1, state_2, q_func, epsilon):
2     """Returns an action selected by an epsilon-Greedy exploration poli
3
4     Args:
5         state_1, state_2 (int, int): two indices describing the current
6         q_func (np.ndarray): current Q-function
7         epsilon (float): the probability of choosing a random command
8
9     Returns:
10        (int, int): the indices describing the action/object to take
11    """
12    action_index, object_index = None, None
13    if np.random.random() < epsilon:
14        action_index, object_index = np.random.randint(0, NUM_ACTIONS),
15    else:

```

Press ESC then TAB or click outside of the code editor to exit

Generating Speech Output

Correct

```
def epsilon_greedy(state_1, state_2, q_func, epsilon):
    """Returns an action selected by an epsilon-Greedy exploration policy

    Args:
        state_1, state_2 (int, int): two indices describing the current state
        q_func (np.ndarray): current Q-function
        epsilon (float): the probability of choosing a random command

    Returns:
        (int, int): the indices describing the action/object to take
    """
    coin = np.random.random_sample()
    if coin < epsilon:
        action_index = np.random.randint(NUM_ACTIONS)
        object_index = np.random.randint(NUM_OBJECTS)
    else:
        q_values = q_func[state_1, state_2, :, :]
        (action_index,
         object_index) = np.unravel_index(np.argmax(q_values, axis=None),
                                         q_values.shape)

    return (action_index, object_index)
```

## Test results

**CORRECT**

[See full output](#)

[See full output](#)

Submit

You have used 4 of 25 attempts

**i** Answers are displayed within the problem

Generating Speech Output


**Topic:** Unit 5 Reinforcement Learning (2 weeks) :Project 5: Text-Based Game / 3. Q-learning Algorithm

Hide Discussion

Add a Post



Show all posts


by recent activity

- 


Test file for tabular\_q\_learning() and epsilon\_greedy()

I've added a test file here: <https://github.com/Praful/MITx-6.86x-version1/tree/master/tests/p...>


 Pinned  Community TA

2
-  What happens when episode ends?


When terminal is True, I understood that the rewards are immediately equal to zero. Q\_value ...

1 new\_
-  How to check these functions?


Is there any test function in the code?

3
-  help staff:tabular\_q\_learning gets right answer for terminal= true,but can not find why terminal= false is wrong


for terminal = false, sequence is right, q\_func is wrong,thanks. q\_func = current\_q + ALPHA \* (...)

5
-  Staff-Info


Hello, I am doing wrong, despite the code seems correct. For the Q function, I update the cur...

2
-  Epsilon-greedy exploration


My understanding of exploration is `action_index = np.random.randint(0, NUM_ACTIONS-1)` o...

4 new\_
-  What is in the Q\_func?


Is the Q function set up as `current room (current state1), current quest(current state 2), actio...`

7
-  How to choose "randomly" the action and object index that matches the grader


To generate the object and action indexes in the the Epsilon-greedy exploration I'm using: `np...`

2
-  What is terminal's role in tabular\_q\_learning function.

I am not sure whether I understand it correctly or not but ,for me, I understand that if we fou...

2
-  Exploitation in Epsilon Greedy function

My ``epsilon_greedy()`` function has code for both exploration and exploitation. The explorati...

1 new\_
-  Stuck on Single Step Update, terminal=True.

I can't seem to get the update right when terminal=True, and I don't know why. I have a condi...

4

Generating Speech Output

