

EdX and its Members use cookies and other tracking technologies for performance, analytics, and marketing purposes. By using this website, you accept this use. Learn more about these technologies in the [Privacy Policy](#).



[Unit 5 Reinforcement Learning](#)(2  
[Course](#) > [weeks](#)) > [Project 5: Text-Based Game](#) >  
8. Linear Q-Learning

**Audit Access Expires May 25, 2020**

You lose all access to this course, including your progress, on May 25, 2020.

## 8. Linear Q-Learning

In this tab, you will implement the Q-learning algorithm with linear function approximation.

Recall the linear approximation we chose.

$$Q(s, c, \theta) = \phi(s, c)^T \theta$$

with

$$\phi(s, c) = \begin{bmatrix} \mathbf{0} \\ \vdots \\ \mathbf{0} \\ \psi_R(s) \\ \mathbf{0} \\ \vdots \\ \mathbf{0} \end{bmatrix}$$

Now, define  $\hat{\theta}_i$  for  $i$  in range  $1, d_C$  so that:

$$\theta = \begin{bmatrix} \hat{\theta}_1 \\ \vdots \\ \hat{\theta}_i \\ \vdots \\ \hat{\theta}_{d_C} \end{bmatrix}$$

With this notation, we get:

$$Q(s, c, \theta) = \psi_R(s)^T \hat{\theta}_c$$

In practice, we can implement  $\hat{\theta}$  as a 2D array, so that

$$\begin{bmatrix} Q(s, 1, \theta) \\ \vdots \\ Q(s, d_C, \theta) \end{bmatrix} = \begin{bmatrix} \hat{\theta}_1^T \\ \vdots \\ \hat{\theta}_{d_C}^T \end{bmatrix} \cdot \psi_R(s)$$

## Epsilon-greedy exploration

1.0/1 point (graded)

Now you will write a function `epsilon_greedy` that implements the  $\epsilon$ -greedy exploration policy using the current Q-function.

**Hint:** You can access  $Q(s, c, \theta)$  using

```
q_value = (theta @ state_vector)[tuple2index(action_index, object_index)]
```

**Available Functions:** You have access to the NumPy python library as `np` and functions `tuple2index` and `index2tuple`. Your code should also use constants `NUM_ACTIONS` and `NUM_OBJECTS`

```
1 def epsilon_greedy(state_vector, theta, epsilon):
2     """Returns an action selected by an epsilon-greedy exploration policy
3
4     Args:
5         state_vector (np.ndarray): extracted vector representation
6         theta (np.ndarray): current weight matrix
7         epsilon (float): the probability of choosing a random command
8
9     Returns:
10        (int, int): the indices describing the action/object to take
11    """
12    if np.random.random() < epsilon:
13        action_index, object_index = np.random.randint(0, NUM_ACTIONS),
14                                     np.random.randint(0, NUM_OBJECTS)
15    else:
```

Press ESC then TAB or click outside of the code editor to exit

Correct

```
def epsilon_greedy(state_vector, theta, epsilon):
    """Returns an action selected by an epsilon-greedy exploration policy

    Args:
        state_vector (np.ndarray): extracted vector representation
        theta (np.ndarray): current weight matrix
        epsilon (float): the probability of choosing a random command

    Returns:
        (int, int): the indices describing the action/object to take
    """
    coin = np.random.random_sample()
    if coin < epsilon:
        action_index = np.random.randint(NUM_ACTIONS)
        object_index = np.random.randint(NUM_OBJECTS)
    else:
        q_values = theta @ state_vector
        index = np.argmax(q_values)
        action_index, object_index = index2tuple(index)
    return (action_index, object_index)
```

## Test results

**CORRECT**

[See full output](#)

[See full output](#)

Submit

You have used 2 of 25 attempts

**i** Answers are displayed within the problem

## Linear Q-learning

1.0/1 point (graded)

Write a function `linear_q_learning` that updates the theta weight matrix, given the transition date  $(s, a, R(s, a), s')$ .

Generating Speech Output

**Reminder:** You should implement this function locally first. You should test this function along with the next one and make sure you achieve reasonable performance

**Hint:** You can access  $Q(s, a, \theta)$  using

```
q_value = (theta @ state_vector)[tuple2index(action_index, object_index)]
```

**Available Functions:** You have access to the NumPy python library as `np`. You should also use constants `ALPHA` and `GAMMA` in your code

```
1 def linear_q_learning(theta, current_state_vector, action_index, object_index,
2                       reward, next_state_vector, terminal):
3     """Update theta for a given transition
4
5     Args:
6         theta (np.ndarray): current weight matrix
7         current_state_vector (np.ndarray): vector representation of current state
8         action_index (int): index of the current action
9         object_index (int): index of the current object
10        reward (float): the immediate reward the agent receives from the environment
11        next_state_vector (np.ndarray): vector representation of next state
12        terminal (bool): True if this episode is over
13
14    Returns:
15        None
```

Press ESC then TAB or click outside of the code editor to exit

Correct

```
def linear_q_learning(theta, current_state_vector, action_index, object_index,
                      reward, next_state_vector, terminal):
    """Update theta for a given transition

    Args:
        theta (np.ndarray): current weight matrix
        current_state_vector (np.ndarray): vector representation of current state
        action_index (int): index of the current action
        object_index (int): index of the current object
        reward (float): the immediate reward the agent receives from playing current action
        next_state_vector (np.ndarray): vector representation of next state
        terminal (bool): True if this episode is over

    Returns:
        None
    """
    q_values_next = theta @ next_state_vector
    maxq_next = np.max(q_values_next)

    q_values = theta @ current_state_vector
    cur_index = tuple2index(action_index, object_index)
    q_value_cur = q_values[cur_index]

    target = reward + GAMMA * maxq_next * (1 - terminal)

    theta[cur_index] = theta[cur_index] + ALPHA * (
        target - q_value_cur) * current_state_vector
```

## Test results

**CORRECT**

[See full output](#)

[See full output](#)

Submit

You have used 15 of 25 attempts

**i** Answers are displayed within the problem

Generating Speech Output

## Evaluate linear Q-learning on Home World game

1/1 point (graded)

Adapt your `run_episode` function to call `linear_Q_learning` and evaluate your performance using hyperparameters:

Set `NUM_RUNS = 5`, `NUM_EPIS_TRAIN = 25`, `NUM_EPIS_TEST = 50`,  $\gamma = 0.5$ , `TRAINING_EP = 0.5`, `TESTING_EP = 0.05` and the **learning rate**  $\alpha = 0.01$ .

Please enter the *average episodic rewards* of your Q-learning algorithm when it converges.

✓ Answer: 0.37

Submit

You have used 2 of 6 attempts

ⓘ Answers are displayed within the problem

## Discussion

Hide Discussion

**Topic:** Unit 5 Reinforcement Learning (2 weeks) :Project 5: Text-Based Game / 8. Linear Q-Learning

Add a Post

Show all posts

by recent activity

💬 [\[Tips from staff:\] Evaluate linear Q-learning on Home World game](#)

2

📌 Pinned   👤 Staff

💬 [updating theta](#)

12 new\_

[a small guidance which helped me: each update should not be for a full theta matrix, but for ...](#)

✓ [What is a dimension of theta and how to understand the hints equation](#)

1 new\_ 5

💬 [plot](#)

5

[There is a plt.show\(\) missing from the end of the agent\\_linear skeleton code](#)

Generating Speech Output

?	<u>[Staff] Solver Problem</u>	2
	Hi, I got this error message "There was a problem running your solution (Staff debug: L379)."...	
?	<u>Epsilon-greedy</u>	2
	I think I'm overthinking the epsilon-greedy implementation... I understand that the q-value is ...	
💬	<u>My converged reward is not liked by the grader..</u>	4
	Community TA	
💬	<u>Using the default hyperparameters won't let you pass the last question</u>	7
💬	<u>linear_q_learning ( theta .... - what is the point to pass it as a parameter if it is global?</u>	6
	Does anybody have and idea?	
✓	<u>How do you find <math>\phi(s, c)</math>?</u>	9
	I'm feeling really dense, but how do you come up with $\phi(s, c)$ for updating theta? Do I need ...	
✓	<u>[Staff] Should theta be returned by linear_q_learning?</u>	2
	Currently, the comment says that linear_q_learning() returns None; It is not clear to me how t...	
✓	<u>[Staff or Student] q_value = (theta @ state_vector)[tuple2index(action_index, object_index)]</u>	3
	Greetings, I'm having a hard time interpreting this for some reason. 1) theta @ state_vector yi...	

© All Rights Reserved

Generating Speech Output